

L103175 - Deploy Ceph Rados Gateway as a replacement for OpenStack Swift

Introduction	1
Openstack Lab Architecture	3
Prepare the environment	4
Check your lab VMs	4
Access the undercloud	5
Create the rados gateway role	6
Configure the rados gateway role	7
OpenStack deployment	9
What next ?	12
Check everything is properly deployed	12
OpenStack	12
Rados Gateway	13
Create a testing environment	15
Openstack	15
Ceph Rados Gateway	16
Swift API	16
S3 API	17
Thank you !	19

Introduction

First of all, it's our pleasure to welcome you to Red Hat Summit 2017, here at the Boston Convention and Exhibition Centre! The past few years have been an exciting time for both Red Hat and the OpenStack community; we've seen unprecedented interest and development in this new revolutionary technology and we're proud to be at the heart of it all.

Red Hat is firmly committed to the future of OpenStack. Our goal is to continue to enhance the technology, make it more readily consumable and to enable our customers to be successful when using it.

This hands-on lab aims to get you understand how you can leverage Red Hat OpenStack Director to deploy Ceph Rados Gateway as a replacement for Openstack Swift.

In the lab you will:

- Learn how to Deploy Ceph Rados Gateway as a replacement for OpenStack Swift using Red Hat OpenStack Director
- Understand how Ceph Rados Gateway is configured to talk to OpenStack Keystone v3 API
- Acquire skills on basic Ceph Rados Gateway commands and interactions
- Gain knowledge on interacting with Ceph Rados Gateway using both Swift and Amazon S3 APIs.

At the end of this session, you will have a good understanding on how to take advantage of Object Storage with a single unified storage backend.

Openstack Lab Architecture

The lab's environment is entirely virtualized and consists of 6 VMs :

Ironi-name	Function	vCPU	vRAM	NIC	Disk
overcloud-ctrl01	Controller	2	8GB	1NIC Provisioning 2NICs data	1x60GB OS
overcloud-ceph01	Ceph OSD	4	4GB	1NIC Provisioning 2NICs data	1x60GB OS 1x60GB OSD
overcloud-ceph02	Ceph OSD	4	4GB	1NIC Provisioning 2NICs data	1x60GB
overcloud-ceph03	Ceph OSD	4	4GB	1NIC Provisioning 2NICs data	1x60GB
overcloud-compute01	Compute	4	8GB	1NIC Provisioning 2NICs data	1x60GB
overcloud-rgw01	RGW	2	4GB	1NIC Provisioning 2NICs data	1x60GB

Plus the Undercloud VM which is obviously not registered into Ironi (4vCPU, 16GB vRAM, 2xNICs (1 data, 1 provisioning), 1x 40GB Disk).

IMPORTANT: You may have noticed that we only have one controller node, this is because we are limited by the lab's hardware. In real life deployments **we strongly recommend to deploy multiple controllers in a HA fashion.**

This lab focuses only on RadosGW implementation, therefore we provide a ready to use undercloud environment and set of templates that configures the following :

- Nodes already registered into ironi
 - Tagged nodes profiles (except for RGW)
- Network isolation, one VLAN per usage (internalAPI, storage, tenant, etc) except for the external network (untagged).
- NIC config that uses the first NIC for provisioning plus a bond on second and third NICs for Openstack traffic.
- Enable Ceph for Nova,Glance and Cinder

Network Name	CIDR	VLAN ID
ControlPlane	172.16.0.0/24	Untagged
InternalAPI	172.17.1.0/24	101
Tenant	172.17.2.0/24	201
Storage	172.17.3.0/24	301
StorageMgmt	172.17.4.0/24	401
External	192.168.122.0/24	Untagged

Prepare the environment

Check your lab VMs

All lab's VMs are tagged with the lab ID (L103175), double check they are all defined and shut-off.

```
hypervisor# virsh list --all | grep L103175
- L103175-undercloud          shut off
- overcloud-L103175-ceph01    shut off
- overcloud-L103175-ceph02    shut off
- overcloud-L103175-ceph03    shut off
- overcloud-L103175-compute01 shut off
- overcloud-L103175-ctrl01    shut off
```

Also check that the provisioning network is active

```
hypervisor# virsh net-list | grep provisioning
provisioning      active      yes         yes
```

Access the undercloud

First you need to boot up the undercloud VM and connect to it.

```
hypervisor# virsh start L103175-undercloud
hypervisor# ssh stack@undercloud
[password : redhat]
undercloud$ source ~/stackrc
```

IMPORTANT : All the steps mentioned below **MUST** be executed as “stack” user.

Note: The ready to use templates mentioned earlier are stored in /home/stack/templates/

List the available ironic nodes and check the maintenance flag **is set to “False”** also verify the provisioning state is “available”

```
undercloud$ openstack baremetal node list --field name power_state provisioning_state
maintenance
```

Name	Power State	Provisioning State	Maintenance
overcloud-ctrl01	power off	available	False
overcloud-compute01	power off	available	False
overcloud-ceph01	power off	available	False
overcloud-ceph02	power off	available	False
overcloud-ceph03	power off	available	False
overcloud-rgw	power off	available	False

If any of your nodes has a maintenance flag set to “True” use this command to turn it off.

```
undercloud$ for node in $(ironic node-list --fields name maintenance | grep True | awk '{print $2}'); do ironic node-set-maintenance $node false; done
```

Create the rados gateway role

As explained during our presentation, we need to create a Rados Gateway role which will be used by director to deploy our RGW node.

Below the list of required steps :

- Define the Rados Gateway Role
- Create the Rados Gateway flavor
- Tag the Ironic node
- Configure the heat templates
 - Role matching and number of nodes
 - Configure network ports
 - Configure NICs attribution.

Copy and edit the default role file to your home directory

```
undercloud$ cp /usr/share/openstack-tripleo-heat-templates/roles_data.yaml
templates/custom_roles.yaml
undercloud$ vi templates/custom_roles.yaml
```

By default the Rados Gateway service is defined into the controller role, as we want a dedicated RGW server, **remove or comment this line** from the Controller role section (line 27):

```
- OS::TripleO::Services::CephRgw # DELETE OR COMMENT
```

Then add the following block at the end of the file, make sure to respect indentation

```
- name: RadosGW
  CountDefault: 1
  HostnameFormatDefault: '%stackname%-rgw-%index%'
  ServicesDefault:
    - OS::TripleO::Services::CACerts
    - OS::TripleO::Services::Kernel
    - OS::TripleO::Services::Ntp
    - OS::TripleO::Services::Timezone
    - OS::TripleO::Services::Snmp
    - OS::TripleO::Services::TripleoPackages
    - OS::TripleO::Services::TripleoFirewall
    - OS::TripleO::Services::SensuClient
    - OS::TripleO::Services::FluentdClient
    - OS::TripleO::Services::CephRgw
    - OS::TripleO::Services::VipHosts
    - OS::TripleO::Services::CephClient
```

That's it ! Your role definition is complete. Now let's configure the node attribution details.

Configure the rados gateway role

Next step is to configure director so **it picks the right Ironic node(s) for this role**. First create and configure a “rgw” nova flavor :

```
undercloud$ openstack flavor create --id auto --ram 4096 --disk 40 --vcpus 1 rgw
undercloud$ openstack flavor set --property "capabilities:boot_option"="local" --property
"capabilities:profile"="rgw" rgw
undercloud$ openstack flavor show rgw | grep properties
| properties                | capabilities:boot_option='local', capabilities:profile='rgw' |
```

Then tag “overcloud-rgw” Ironic node with the profile we just defined above. By default all our nodes except overcloud-rgw has a profile assigned, let’s assign it one !

```
undercloud$ openstack overcloud profiles list
```

Node UUID	Node Name	Provision State	Current Profile	Possible Profiles
b2c14304-7c0f-42a2-9546-74fb2d950776	overcloud-ctrl01	available	control	
6821b53a-1c9c-409d-8b7d-6d252a74e32e	overcloud-compute01	available	compute	
86ba67ce-daa6-4128-abe8-0e05c8d1fc10	overcloud-ceph01	available	ceph-storage	
fac3d6bc-61a4-4546-afbc-56d4f38b3ba0	overcloud-ceph02	available	ceph-storage	
56096f50-1b1f-4279-8818-43a432d77ab5	overcloud-ceph03	available	ceph-storage	
e2442bbc-0468-477b-81c2-cf99d8be3d62	overcloud-rgw	available	None	

Tag overcloud-rgw with the rgw profile

```
undercloud$ ironic node-update overcloud-rgw add
properties/capabilities='profile:rgw,boot_option:local'
```

```
undercloud$ openstack baremetal node show overcloud-rgw | grep properties
| properties                | {'memory_mb': u'4096', 'cpu_arch': u'x86_64', 'local_gb': u'59',
u'cpus': u'2', u'capabilities': u'profile:rgw,boot_option:local'}
```

```
undercloud$ openstack overcloud profiles list
```

Node UUID	Node Name	Provision State	Current Profile	Possible Profiles
b2c14304-7c0f-42a2-9546-74fb2d950776	overcloud-ctrl01	active	control	
6821b53a-1c9c-409d-8b7d-6d252a74e32e	overcloud-compute01	active	compute	
86ba67ce-daa6-4128-abe8-0e05c8d1fc10	overcloud-ceph01	active	ceph-storage	
fac3d6bc-61a4-4546-afbc-56d4f38b3ba0	overcloud-ceph02	active	ceph-storage	
56096f50-1b1f-4279-8818-43a432d77ab5	overcloud-ceph03	active	ceph-storage	
e2442bbc-0468-477b-81c2-cf99d8be3d62	overcloud-rgw	active	rgw	

Last step is to tell **director to pick a node that has the capability “rgw”** whenever a RGW role node needs to be deployed; we take this opportunity to configure the number of RGW node(s) we’d like to deploy, in our case just one.

```
undercloud$ cat << EOF > ~/templates/rgw-flavor.yaml
parameter_defaults:
  OvercloudRadosGWFlavor: rgw
  RadosGWCount: 1
EOF
```

Note: Setting RadosGWCount to 1 is optional here as we already have CountDefault: 1 in the role file. This is only meant to demonstrate how to override the default value.

Our new role definition is almost complete, the only part missing is **the network configuration**. We need to tell director on **which subnets this role needs to be connected and how to configure the node’s NICs**.

Create a rgw-network.yaml file:

```
undercloud$ cat << EOF > ~/templates/rgw-network.yaml
resource_registry:
  OS::TripleO::RadosGW::Ports::ExternalPort:
    /usr/share/openstack-tripleo-heat-templates/network/ports/noop.yaml
  OS::TripleO::RadosGW::Ports::InternalApiPort:
    /usr/share/openstack-tripleo-heat-templates/network/ports/internal_api.yaml
  OS::TripleO::RadosGW::Ports::StoragePort:
    /usr/share/openstack-tripleo-heat-templates/network/ports/storage.yaml
  OS::TripleO::RadosGW::Ports::StorageMgmtPort:
    /usr/share/openstack-tripleo-heat-templates/network/ports/noop.yaml
  OS::TripleO::RadosGW::Ports::TenantPort:
    /usr/share/openstack-tripleo-heat-templates/network/ports/noop.yaml

  OS::TripleO::RadosGW::Net::SoftwareConfig: /home/stack/templates/nic-configs/rgw.yaml
EOF
```

We ask director to connect our Rados Gateway Role (i.e nodes) to the following networks :

- Internal API
- Storage Access

Note: Public Access to RGW is handled by HAProxy which by default is hosted on the controllers therefore there is no need to have a connection to the external network.

Then the last line tells Director to use “/home/stack/templates/nic-configs/rgw.yaml” **to configure the NICs**.

There is no default pre-built nic config template shipped for RGW, for simplicity sake, we're going to use the swift-storage template and adjust it to our needs.

The only required change is to **remove the storage_mgmt interface** which is required for Swift but not for RGW.

Delete the appropriate lines so our RGW is not connected to the storage_mgmt network.

```
undercloud$ cd ~/templates/nic-configs/  
undercloud$ sed -e '150,156d' swift-storage.yaml > rgw.yaml
```

If you're curious about which lines have been deleted, you can use diff to satisfy your curiosity

```
undercloud$ diff swift-storage.yaml rgw.yaml  
150,156d149  
<             -  
<                 type: vlan  
<                 device: bond1  
<                 vlan_id: {get_param: StorageMgmtNetworkVlanID}  
<                 addresses:  
<                     -  
<                     ip_netmask: {get_param: StorageMgmtIpSubnet}
```

Note: *When doing such things in real life, we recommend to also update the description section at the beginning of the template file.*

OpenStack deployment

Our configuration is now complete, time to kick off the deployment, we provided a deployment script (~stack/templates/overcloud-deploy.sh) which **we need to modify to reflect the RGW addition**.

Edit ~stack/templates/overcloud-deploy.sh and **add the following lines at the end of the file** :

```
-e ~/templates/rgw-network.yaml \  
-e ~/templates/rgw-flavor.yaml \  
-e /usr/share/openstack-tripleo-heat-templates/environments/ceph-radosgw.yaml \  
-r ~/templates/custom_roles.yaml
```

We include our two RGW templates to the environment (rgw-network.yaml and rgw-flavor.yaml) as well as the custom role file we created earlier.

Swift being the default object storage backend in Red Hat OSP 10, RGW is de facto disabled (set to OS::Heat::None). **The ceph-radosgw.yaml template enables RGW and disable all Swift services**, inclusion of this template is mandatory to use RGW.

Your deployment script should now looks like this - **double check it is the same as below** :

```
undercloud$ cat ~/templates/overcloud-deploy.sh
#!/bin/bash

exec openstack overcloud deploy \
    --templates --ntp-server 10.16.255.1 \
    --control-flavor control --control-scale 1 \
    --compute-flavor compute --compute-scale 1 \
    --ceph-storage-flavor ceph-storage --ceph-storage-scale 3 \
    --neutron-tunnel-types vxlan --neutron-network-type vxlan \
    -e /usr/share/openstack-tripleo-heat-templates/environments/network-isolation.yaml \
    -e ~/templates/storage-environment.yaml \
    -e ~/templates/network-environment.yaml \
    -e ~/templates/rgw-network.yaml \
    -e ~/templates/rgw-flavor.yaml \
    -e /usr/share/openstack-tripleo-heat-templates/environments/ceph-radosgw.yaml \
    -r ~/templates/custom_roles.yaml
```

Kick off the deployment :

```
undercloud$ screen
undercloud$ source ~/stackrc
undercloud$ cd ~
undercloud$ sh templates/overcloud-deploy.sh
Started Mistral Workflow. Execution ID: xxxx
Plan updated
Deploying templates in the directory /tmp/tripleoclient-xxxx/tripleo-heat-templates
```

This process takes some time, while you wait for the deployment to complete you can have a look at the ironic nodes being booted up (create a new “screen” window or use another terminal).

```
undercloud$ source ~/stackrc
undercloud$ openstack baremetal node list
```

UUID	Name	Instance UUID	Power State	Provisioning State	Maintenance
b2c14304-7c0f-42a2-9546-74fb2d950776	overcloud-ctrl01	6719079a-bf30-4def-b1da-5f179c876a63	power on	deploying	False
6821b53a-1c9c-409d-8b7d-6d252a74e32e	overcloud-compute01	fb15a764-9f14-4772-aba1-0cba96bdf6af	power on	deploying	False
86ba67ce-daa6-4128-abe8-0e05c8d1fc10	overcloud-ceph01	21859eb5-aea6-4ff0-9853-595914326a78	power on	wait call-back	False

fac3d6bc-61a4-4546-afbc-56d4f38b3ba0	overcloud-ceph02	3c006ffd-7168-4815-8b3e-7c7c94a44724	power on	deploying	False	
56096f50-1b1f-4279-8818-43a432d77ab5	overcloud-ceph03	8d848c9d-bd3d-4639-89db-adc25faef9fa	power on	deploying	False	
e2442bbc-0468-477b-81c2-cf99d8be3d62	overcloud-rgw	b327de8f-e8a2-4dac-8bf3-751fb2e13f7b	power on	deploying	False	

Also have a look at the nova nodes being booted up

undercloud\$ openstack server list

ID	Name	Status	Networks	Image Name
21859eb5-aea6-4ff0-9853-595914326a78	overcloud-cephstorage-2	BUILD	ctlplane=172.16.0.26	overcloud-full
8d848c9d-bd3d-4639-89db-adc25faef9fa	overcloud-cephstorage-0	BUILD	ctlplane=172.16.0.23	overcloud-full
6719079a-bf30-4def-b1da-5f179c876a63	overcloud-controller-0	BUILD	ctlplane=172.16.0.27	overcloud-full
3c006ffd-7168-4815-8b3e-7c7c94a44724	overcloud-cephstorage-1	BUILD	ctlplane=172.16.0.22	overcloud-full
b327de8f-e8a2-4dac-8bf3-751fb2e13f7b	overcloud-rgw-0	BUILD	ctlplane=172.16.0.29	overcloud-full
fb15a764-9f14-4772-aba1-0cba96bdf6af	overcloud-compute-0	BUILD	ctlplane=172.16.0.28	overcloud-full

When the deployment finishes, you should see something like :

Stack overcloud CREATE_COMPLETE

Started Mistral Workflow. Execution ID: xxxxx

/home/stack/.ssh/known_hosts updated.

Original contents retained as /home/stack/.ssh/known_hosts.old

Overcloud Endpoint: http://192.168.122.xxx:5000/v2.0

Overcloud Deployed

Check all nodes are up and deployed :

undercloud\$ openstack baremetal node list

UUID	Name	Instance UUID	Power State	Provisioning State	Maintenance
b2c14304-7c0f-42a2-9546-74fb2d950776	overcloud-ctrl01	6719079a-bf30-4def-b1da-5f179c876a63	power on	active	False
6821b53a-1c9c-409d-8b7d-6d252a74e32e	overcloud-compute01	fb15a764-9f14-4772-aba1-0cba96bdf6af	power on	active	False
86ba67ce-daa6-4128-abe8-0e05c8d1fc10	overcloud-ceph01	21859eb5-aea6-4ff0-9853-595914326a78	power on	active	False
fac3d6bc-61a4-4546-afbc-56d4f38b3ba0	overcloud-ceph02	3c006ffd-7168-4815-8b3e-7c7c94a44724	power on	active	False
56096f50-1b1f-4279-8818-43a432d77ab5	overcloud-ceph03	8d848c9d-bd3d-4639-89db-adc25faef9fa	power on	active	False
e2442bbc-0468-477b-81c2-cf99d8be3d62	overcloud-rgw	b327de8f-e8a2-4dac-8bf3-751fb2e13f7b	power on	active	False

undercloud\$ openstack server list

ID	Name	Status	Networks	Image Name
21859eb5-aea6-4ff0-9853-595914326a78	overcloud-cephstorage-2	ACTIVE	ctlplane=172.16.0.26	overcloud-full
8d848c9d-bd3d-4639-89db-adc25faef9fa	overcloud-cephstorage-0	ACTIVE	ctlplane=172.16.0.23	overcloud-full
6719079a-bf30-4def-b1da-5f179c876a63	overcloud-controller-0	ACTIVE	ctlplane=172.16.0.27	overcloud-full
3c006ffd-7168-4815-8b3e-7c7c94a44724	overcloud-cephstorage-1	ACTIVE	ctlplane=172.16.0.22	overcloud-full
b327de8f-e8a2-4dac-8bf3-751fb2e13f7b	overcloud-rgw-0	ACTIVE	ctlplane=172.16.0.29	overcloud-full

```
| fb15a764-9f14-4772-aba1-0cba96bdf6af | overcloud-compute-0 | ACTIVE | ctlplane=172.16.0.28 | overcloud-full |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

For simplicity sake, populate the undercloud hosts file so we can access the nodes by their names.

```
undercloud$ nova list --fields name,networks | awk '/overcloud/ { gsub("ctlplane=", ""); print
$6" "$4; }' | sudo tee -a /etc/hosts
```

Ping one of the overcloud's node for example the controller

```
undercloud$ ping overcloud-controller-0
PING overcloud-controller-0 (172.16.0.x) 56(84) bytes of data.
64 bytes from overcloud-controller-0 (172.16.0.x): icmp_seq=1 ttl=64 time=0.232 ms
```

What next ?

Check everything is properly deployed

OpenStack

Let's now verify that our overcloud is operational, first we are going to quickly check the core OpenStack services then focus on RGW.

List all running services :

```
undercloud$ source ~/overcloudrc
undercloud$ openstack service list
```

ID	Name	Type
21456c69580640de9a14b573c7de9a59	heat-cfn	cloudformation
5f4f1c37a83541f59be3432cf4b172a4	neutron	network
623e4395e4854353872999fcaf8d8017	nova	compute
78a0809fdf6f4353ac2430a5f599a012	cinder	volume
8875fe6e065c477db787c2cd53f7362f	swift	object-store
91777fc71ce949d49bdd3cd27a5ab4f	cinderv2	volumev2
978f404f32fc496fabfe3fdd8db8cb6f	glance	image
bd63c0a7bca747f695772714a240f8ee	gnocchi	metric
c52bc9dbff174da0b07f41e76fc8b8f9	ceilometer	metering
c61af82e7613463d86bc0fcd43dda666	cinderv3	volumev3
c72b6bf38e5647c8b86b48ab1ab78346	keystone	identity
cd53b7a1262e48a39dcce221480d6188	aodh	alarming

```
| d22e8b9d1c0e4efe9e8d7761322172fa | heat      | orchestration |
+-----+-----+-----+
```

List all OpenStack endpoints :

```
undercloud$ openstack catalog list
```

```
+-----+-----+-----+
-----+
| Name      | Type      | Endpoints
|
+-----+-----+-----+
-----+
| nova      | compute   | regionOne
|
|           |           | publicURL: http://192.168.122.107:8774/v2.1
|
|           |           | internalURL: http://172.17.1.13:8774/v2.1
|
|           |           | adminURL: http://172.17.1.13:8774/v2.1
|
```

(...)

Rados Gateway

In this section we will have a look at how Rados Gateway has been deployed.
Connect to the controller.

```
undercloud$ ssh heat-admin@overcloud-controller-0
```

In a traditional production deployment, HAProxy configuration balances VIPs 8080/tcp requests to all controller nodes on 8080/tcp internal IPs. In our case we only have one controller.

```
overcloud-controller-0$ sudo -i
overcloud-controller-0# grep "listen ceph_rgw" /etc/haproxy/haproxy.cfg -A3
listen ceph_rgw
    bind 172.17.3.13:8080 transparent
    bind 192.168.122.109:8080 transparent
    server overcloud-rgw-0.storage.localdomain 172.17.3.15:8080 check fall 5 inter 2000 rise 2
```

We can confirm that with :

```
overcloud-controller-0# ss -lnp | grep ":8080"
tcp    LISTEN    0      128      192.168.122.109:8080          *:*
users: (("haproxy",pid=149466,fd=10))
```

```
tcp    LISTEN      0      128      172.17.3.13:8080      *:*\nusers:(("haproxy",pid=149466,fd=9))
```

Log Out and connect to the rados gateway node:

```
undercloud$ ssh heat-admin@overcloud-rgw-0
```

Have a look at the RGW process :

```
overcloud-rgw-0$ sudo -i\novercloud-rgw-0# ps aux | grep radosgw\nceph      37012  0.0  0.7 2742588 30540 ?        Ssl  Apr13   1:09 /usr/bin/radosgw -f\n--cluster ceph --name client.radosgw.gateway --setuser ceph --setgroup ceph
```

We have RGW running with client.radosgw.gateway as auth key :

```
overcloud-rgw-0# cat /etc/ceph/ceph.client.radosgw.gateway.keyring\n[client.radosgw.gateway]\n    key = AQCHTvtXAAAAABAAP9TMC7mc7HL/ovNmtqojZA==\n    caps mon = "allow *"\n    caps osd = "allow *"
```

RGW related configuration in /etc/ceph/ceph.conf :

```
[client.radosgw.gateway]\nuser = apache\nrgw_frontends = civetweb port=172.17.3.15:8080\nlog_file = /var/log/ceph/radosgw.log\nhost = overcloud-rgw-0\nkeyring = /etc/ceph/ceph.client.radosgw.gateway.keyring\nrgw_keystone_token_cache_size = 500\nrgw_keystone_url = http://172.16.0.24:35357\nrgw_s3_auth_use_keystone = True\nrgw_keystone_admin_token = QD3Gw4dWGDUZuV92f7Rm2ZABz\nrgw_keystone_accepted_roles = admin,_member_,Member
```

Create a testing environment

Openstack

The best way to quickly check if our Openstack and RGW are operational is to deploy a testing environment.

We start by importing our SSH public key into nova :

```
undercloud$ source ~/overcloudrc
undercloud$ openstack keypair create --public-key ~/.ssh/id_rsa.pub rhsummit
undercloud$ openstack keypair list
```

Name	Fingerprint
rhsummit	20:da:1e:c4:8f:ba:01:b4:89:0c:72:cb:7c:e6:de:86

Create the “admin stack”, this stack creates, networks, flavor, image, user, etc

```
undercloud$ openstack stack create -t ~/heat/admin.yaml stack_admin --wait
```

Once the stack has successfully terminated, list the newly created subnets.

```
undercloud$ openstack network list
```

ID	Name	Subnets
d29b743c-3a20-47e0-a70a-05abe573f5d1	external	fa72ad95-78b2-4dbf-a337-056b7fab6016
d5656f5f-d492-4be3-9268-fb9f1e33d73c	net1	be7a34fb-3887-4512-ba29-eadef75008d9

Here we have

- One external subnet (no tunneling / used for floating IPs)
- One tenant network - “net1” (VXLAN tunneling / private instance network)

Check our flavor type has been created

```
undercloud$ openstack flavor list
```

ID	Name	RAM	Disk	Ephemeral	VCPUs	Is Public
88d59257-9a42-48d2-95e9-4608b77814ac	m1.small	1024	0	1	1	True

Create the “user stack”, this stack creates, one instance and attach a floating IP to it.

```
undercloud$ openstack stack create -t ~/heat/user.yaml stack_user --wait
```

Note: In order to reduce the image size of the lab we use a qcow2 image which slows down the instance boot process. On live environments you should use raw images if you use Ceph as backend.

```
undercloud$ openstack server list
```

ID	Name	Status	Networks	Image Name
4f91fbc9-e4c7-429f-b1f3-96871229b9a5	vm1	ACTIVE	net1=192.168.1.20, 192.168.122.161	centos7

Test the instance's connectivity by pinging its floating IP (192.168.122.x)

```
undercloud$ ping 192.168.122.x
```

Ceph Rados Gateway

In this section we will show you how to use RadosGW from an end-user perspective. We will cover both Swift and S3 APIs.

Swift API

First take note of your overcloud external API URL, you will need it in few moments

```
undercloud$ grep OS_AUTH_URL overcloudrc
export OS_AUTH_URL=http://192.168.122.xxx:5000/v2.0
```

Next connect to the instance :

```
undercloud$ source ~/overcloudrc
```

```
undercloud$ openstack server list
```

ID	Name	Status	Networks	Image Name
4f91fbc9-e4c7-429f-b1f3-96871229b9a5	vm1	ACTIVE	net1=192.168.1.20, 192.168.122.161	centos7

```
undercloud$ ssh centos@192.168.122.x
```


Get root privileges and jump into the swift directory and edit the userrc file:

```
vm1$ sudo -i
vm1$ cd swift
vm1$ vi userrc
```

And replace KEYSTONE_ENDPOINT_ADDRESS with the IP address of the Keystone public endpoint. Look at the beginning of this section on how to find it.

Then execute the swift-exercise.sh script:

```
vm1$ ./swift-exercise.sh
```

This will execute a bunch of commands against the Ceph Rados Gateway using the Swift API. Have a look at the script, it contains commands along with comments on what does what !

Now, logout from your guest virtual machine:

```
vm1$ logout
```

S3 API

Since we want to test the S3 API from Ceph Rados Gateway, we need to create a user on it, for this log on the Ceph Rados Gateway machine and create the user:

```
undercloud$ ssh heat-admin@overcloud-rgw-0
```

```
overcloud-rgw-0$ sudo -i
```

```
overcloud-rgw-0# radosgw-admin user create --uid=user1 --display-name="Ceph demo user"
--access-key=G1EZ5R4K6IJ7XUQKMAED --secret-key=cNmUrqpBKjCMzcfqG8fg4Qk07Xkoyau520mvmSsz
```

```
{
  "user_id": "user1",
  "display_name": "Ceph demo user",
  "email": "",
  "suspended": 0,
  "max_buckets": 1000,
  "aud": 0,
  "subusers": [],
  "keys": [
    {
      "user": "user1",
      "access_key": "G1EZ5R4K6IJ7XUQKMAED",
      "secret_key": "cNmUrqpBKjCMzcfqG8fg4Qk07Xkoyau520mvmSsz"
```

```

    }
  ],
  "swift_keys": [],
  "caps": [],
  "op_mask": "read, write, delete",
  "default_placement": "",
  "placement_tags": [],
  "bucket_quota": {
    "enabled": false,
    "max_size_kb": -1,
    "max_objects": -1
  },
  "user_quota": {
    "enabled": false,
    "max_size_kb": -1,
    "max_objects": -1
  },
  "temp_url_keys": []
}

```

Logout back to your undercloud and take note of the address of your Swift API, using your overcloudrc credentials:

```

undercloud$ source ~/overcloudrc
undercloud$ openstack catalog list | grep swift | grep public
publicURL: http://192.168.122.109:8080/swift/v1

```

Go back into your guest virtual machine again and become root:

```

undercloud$ ssh centos@192.168.122.x
vm1$ sudo -i

```

Let's test the S3 API now, edit the .s3cfg file

```
vm1# vi .s3cfg
```

Change these two options with the Swift publicURL we just got above:

```

host_base = SWIFT_ENDPOINT_ADDRESS:8080
host_bucket = SWIFT_ENDPOINT_ADDRESS:8080

```

Now, change directory and execute the s3-exercise.sh :

```

vm1# cd s3/
vm1# ./s3-exercise.sh

```

Have a look at the script, it contains commands along with comments on what does what !

Thank you !

That's all ! Thanks a lot for attending our session!

We hope it met your expectations and you learnt how you can use Red Hat OpenStack director to leverage Red Hat Ceph Storage as a replacement for Swift!

Don't hesitate to come meet us and ask anything.

Have a great day and enjoy summit !

Sebastien Han - Red Hat Ceph Storage Engineering

Gregory Charot - Red Hat Openstack Platform Field Product Management

Cyril Lopez - Red Hat Consulting

<https://www.redhat.com/en/topics/openstack>

<https://www.redhat.com/en/technologies/linux-platforms/openstack-platform>

<https://www.redhat.com/en/technologies/storage/ceph>